

**CINECA**

HIGH PERFORMANCE  
COMPUTING CINECA  
ITALY

# Parallel Programming with MPI

## Part III - Derived datatypes



HPC SCHOOL  
— COSTA RICA —

**CINECA**



**ALESSANDRO MARANI**

a.marani@ cineca.it || San José 2023

# AGENDA

## DERIVED DATATYPES

Definition, typemap, how to use

---

## DERIVED DATATYPE CONSTRUCTORS

Contiguous, vector and other examples

---

## MANAGING STRUCTS WITH CONSTRUCTORS

MPI\_Type\_create\_struct, extent, using displacements

---



A blue-tinted photograph of a large banner for the Leonardo project. The banner features the Leonardo logo, the European Union flag, the text 'Funded by the European Union', and the acronym 'CINECA'. The banner is set against a background of server racks in a data center.

# Derived datatypes



# Need for derived datatypes


**You may need to send messages that contain:**

1. **non-contiguous data** of a single type (e.g. a sub-block of a matrix)
2. contiguous data of **mixed types** (e.g., an integer count, followed by a sequence of real numbers)
3. **non-contiguous data** of **mixed types**

**Possible solution:**

Make multiple MPI calls to send and receive each data element

→ If advantageous, copy data to a buffer before sending it



Additional latency costs due to multiple calls  
Additional latency costs due to memory copy


# Datatype Solution

The idea of MPI derived datatypes is to provide a simple, portable, elegant and efficient way of communicating non-contiguous or mixed types in a message.

During the communication, the datatype tells MPI system where to take the data when sending or where to put data when receiving.

**The actual performances depend on the MPI implementation**

Derived datatypes are also needed for getting the most out of MPI-I/O.



# Definitions

A **general datatype** is an opaque object able to describe a buffer layout in memory by specifying:

- A sequence of basic datatypes
- A sequence of integer (byte) displacements.

**Typemap** = {(type 0, displ 0), ... (type n-1, displ n-1)}

- pairs of basic types and displacements (in bytes)

**Type signature** = {type 0, type 1, ... type n-1}

- list of types in the typemap
- gives size of each elements and tells MPI how to interpret the bits it sends and receives

**Displacement:**

- tells MPI where to get (when sending) or put (when receiving)

# Typemap


## Example:

Basic datatype are particular cases of a general datatype, and are predefined:

**MPI\_INT = {(int, 0)}**

General datatype with typemap:

**Typemap = {(int,0), (double,8), (char,16)}**



## How to use

General datatypes are created (and destroyed) at run-time through calls to MPI library routines

### Implementation steps are:

1. Creation of the datatype from existing ones with a **datatype constructor**;
2. Allocation (**committing**) of the datatype before using it;
3. **Usage of the derived datatype** for MPI communications and/or for MPI-I/O;
4. Deallocation (**freeing**) of the datatype after that it is no longer needed;



## Committing and freeing

```
int MPI_Type_commit(MPI_Datatype *type);
```

- ❑ Before it can be used in a communication or I/O call, each derived datatype has to be committed
- ❑ New datatypes have to be declared as MPI\_Datatype

```
int MPI_Type_free(MPI_Datatype *type);
```

- ❑ Mark a datatype for deallocation
- ❑ Datatype will be deallocated when all pending operations are finished

A blue-tinted photograph of a large banner for the Leonardo project. The banner features the word 'LEONARDO' in large, bold, white letters. To the left of the word is a circular logo with a stylized 'L'. To the right of the word are several smaller logos: the European Union flag, the text 'Funded by the European Union', and the text 'CINECA'. The banner is set against a background of a large, modern building with a glass facade.

# Derived datatype constructors

## MPI\_Type\_contiguous

```
int MPI_Type_contiguous(int count, MPI_Datatype oldtype,  
                        MPI_Datatype *newtype);
```

<b>count</b>	number of consecutive elements of the old datatype
<b>oldtype</b>	the previous datatype to be used as base
<b>newtype</b>	the new derived datatype

**MPI\_TYPE\_CONTIGUOUS** constructs a typemap consisting of the replication of a datatype into contiguous locations.

**newtype** is the datatype obtained by concatenating **count** copies of **oldtype**.

# MPI\_Type\_contiguous

```
count = 4;  
MPI_Type_contiguous(count, MPI_FLOAT, &rowtype);
```

1.0	2.0	3.0	4.0
5.0	6.0	7.0	8.0
9.0	10.0	11.0	12.0
13.0	14.0	15.0	16.0

$a[4][4]$

```
MPI_Send(&a[2][0], 1, rowtype, dest, tag, comm);
```

9.0	10.0	11.0	12.0
-----	------	------	------

1 element of  
rowtype



## MPI\_Type\_vector

```
int MPI_Type_vector(int count, int blocklength, int stride,  
                    MPI_Datatype oldtype, MPI_Datatype *newtype);
```

<b>count</b>	number of blocks
<b>blocklength</b>	number of elements in each block
<b>Stride</b>	number of elements (NOT bytes) between the start of each block
<b>oldtype</b>	the previous datatype to be used as base
<b>newtype</b>	the new derived datatype

Consists of a number of elements of the same datatype  
repeated with a certain stride

# MPI\_Type\_vector

```
count = 4;  blocklength = 1;  stride = 4;  
MPI_Type_vector(count, blocklength, stride, MPI_FLOAT,  
                &columntype);
```

1.0	2.0	3.0	4.0
5.0	6.0	7.0	8.0
9.0	10.0	11.0	12.0
13.0	14.0	15.0	16.0

a[4][4]

```
MPI_Send(&a[0][1], 1, columntype, dest, tag, comm);
```

2.0	6.0	10.0	14.0
-----	-----	------	------


1 element of  
columntype

# MPI\_Type\_indexed

```
int MPI_Type_indexed(int count, int *array_of_blocklengths,  
                    int *array_of_displacements, MPI_Datatype oldtype,  
                    MPI_Datatype *newtype);
```

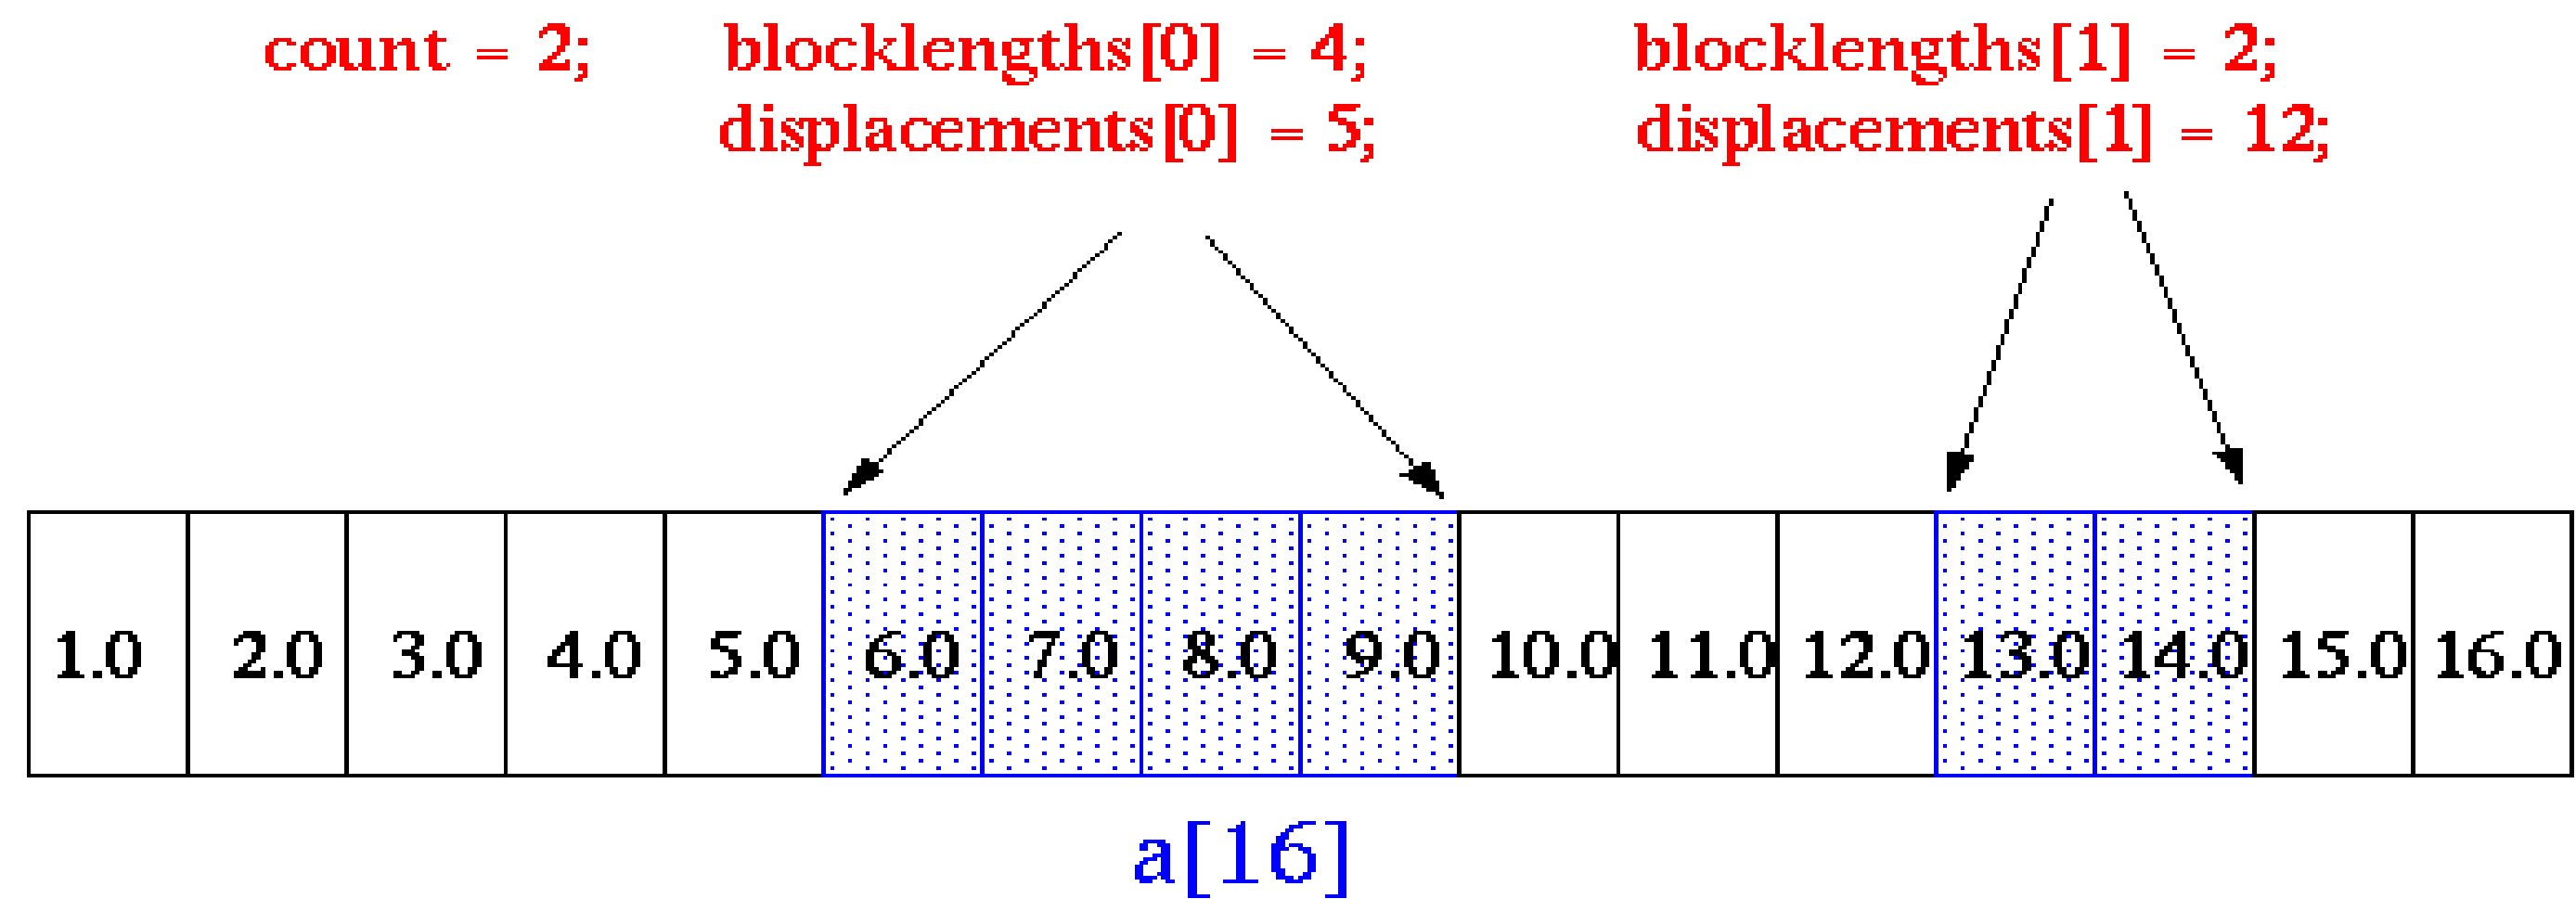
<b>count</b>	number of blocks
<b>array_of_blocklengths</b>	number of elements in each block
<b>array_of_displacements</b>	displacement for each block
<b>oldtype</b>	the previous datatype to be used as base
<b>newtype</b>	the new derived datatype

Creates a new type from blocks comprising identical elements  
The size and displacements of the blocks can vary




count=3, array\_of\_blocklengths=(/2,3,1/), array\_of\_displacements=(/0,3,8/)

# MPI\_Type\_indexed



```
MPI_Type_indexed(count, blocklengths, displacements, MPI_FLOAT, &indextype);
```

```
MPI_Send(&a, 1, indextype, dest, tag, comm);
```





# MPI\_Type\_indexed example


```
/* Lower Triangular Matrix (LTM) */
MPI_Datatype lower;
float *a;
int displ[100], blocklen[100];
a = malloc(100*100*sizeof(float));

/* Compute start and size of the rows */
for (int i=0, i < 100; i++) {
    displ[i]=100*(i-1);
    blocklen[i]=i;
}

/* Create and commit a datatype for the LTM */
MPI_Type_indexed(100, blocklen, disp, MPI_FLOAT, &lower);
MPI_Type_commit(&lower);

/* Use it... */
MPI_Send(a, 1, lower, dest, tag, MPI_COMM_WORLD);

/* Free it after use */
MPI_Type_free(&lower);
```



## MPI\_Type\_subarray

```
int MPI_Type_subarray(int ndims, int *array_of_sizes,  
    int *array_of_subsizes, int *array_of_starts, int order,  
    MPI_Datatype oldtype, MPI_Datatype *newtype);
```

<b>ndims</b>	number of array dimensions
<b>array_of_sizes</b>	number of elements of type “oldtype” in each dimension of the full array
<b>array_of_subsizes</b>	number of elements of type “oldtype” in each dimension of the subarray
<b>array of starts</b>	starting coordinates of the subarray in each dimension
<b>order</b>	array storage order flag (MPI_ORDER_C or MPI_ORDER_FORTRAN)
<b>oldtype</b>	the previous datatype to be used as base
<b>newtype</b>	the new derived datatype

The subarray type constructor creates an MPI datatype describing an  $n$ -dimensional subarray of an  $n$ -dimensional array.

The subarray may be situated anywhere within the full array, and may be of any nonzero size up to the size of the larger array as long as it is confined within this array.

# MPI\_Type\_subarray example



```
double matrix[100][100];
double subarray[100][25];
MPI_Datatype filetype;
int sizes[2], subsizes[2], starts[2];
int rank;

MPI_Comm_rank(MPI_COMM_WORLD, &rank);

sizes[0]=100; sizes[1]=100;
subsizes[0]=100; subsizes[1]=25;
starts[0]=0; starts[1]=rank*subsizes[1];

MPI_Type_create_subarray(2, sizes, subsizes, starts,
MPI_ORDER_C, MPI_DOUBLE, &filetype);

MPI_Type_commit(&filetype);
MPI_Type_send(matrix, 1, filetype, dest, tag, MPI_COMM_WORLD);
MPI_Type_free(&filetype);
```





# Managing structs with constructors



# MPI\_Type\_struct

```
int MPI_Type_struct(int count, int *array_of_blocklengths,  
    int *array_of_displacements, MPI_Datatype array_of_oldtypes,  
    MPI_Datatype *newtype);
```

<b>count</b>	number of blocks
<b>array_of_blocklengths</b>	number of elements in each block
<b>array_of_displacements</b>	BYTE displacement for each block
<b>oldtype</b>	type of elements for each block
<b>newtype</b>	the new derived datatype

- ❑ This subroutine returns a new datatype that represents *count* blocks. Each block is defined by an entry in *array\_of\_blocklengths*, *array\_of\_displacements* and *array\_of\_oldtypes*.
- ❑ Displacements are expressed in **bytes** (since the type can change!)
- ❑ To gather a mix of different datatypes scattered at many locations in space into one datatype that can be used for the communication.


# Using extent

## Extent:

The extent of a datatype is the span from the lower to the upper bound in bytes (including “holes”)

```
struct {  
    float x, y, z, velocity;  
    int n, type;  
} Particle;  
  
Particle particles[NELEM];
```

```
MPI_Type_extent(MPI_FLOAT, &extent);  
count = 2;  
blockcounts[0] = 4;    blockcount[1] = 2;  
oldtypes[0] = MPI_FLOAT; oldtypes[1] = MPI_INT;  
displ[0] = 0;          displ[1] = 4*extent;
```



```
MPI_Type_struct (count, blockcounts, displ, oldtypes, &particletype);  
MPI_Type_commit (&particletype);
```

# Using extent

```
struct {  
    float x, y, z, velocity;  
    int n, type;  
} Particle;  
  
Particle particles[NELEM];
```

```
int count, blockcounts[2];  
MPI_Aint displ[2];  
MPI_Datatype particletype, oldtypes[2];  
  
count = 2;  
blockcounts[0] = 4; blockcounts[1] = 2;  
oldtypes[0] = MPI_FLOAT; oldtypes[1] = MPI_INT;  
  
MPI_Type_extent(MPI_FLOAT, &extent);  
displ[0] = 0; displ[1] = 4*extent;  
  
MPI_Type_create_struct (count, blockcounts, displ, oldtypes,  
                        &particletype);  
MPI_Type_commit(&particletype);  
  
MPI_Send (particles, NELEM, particletype, dest, tag,  
           MPI_COMM_WORLD);  
  
MPI_Type_free(&particletype);
```

# Using displacements

```
struct PartStruct {  
    char class;  
    double d[6];  
    int b[7];  
} particle[100];
```

```
MPI_Datatype ParticleType;  
int count = 3;  
MPI_Datatype type[3] = {MPI_CHAR, MPI_DOUBLE, MPI_INT};  
int blocklen[3] = {1, 6, 7};  
MPI_Aint disp[3];  
  
MPI_Get_address(&particle[0].class, &disp[0]);  
MPI_Get_address(&particle[0].d, &disp[1]);  
MPI_Get_address(&particle[0].b, &disp[2]);  
  
/* Make displacements relative */  
disp[2] -= disp[0]; disp[1] -= disp[0]; disp[0] = 0;  
  
MPI_Type_create_struct (count, blocklen, disp, type, &ParticleType);  
MPI_Type_commit (&ParticleType);  
  
MPI_Send(particle, 100, ParticleType, dest, tag, comm);  
MPI_Type_free (&ParticleType);
```



# Performance

Performance depends on the datatype – more general datatypes are often slower.

Some MPI implementations can handle important special cases: e.g., constant stride, contiguous structures.

Overhead is potentially reduced by sending one long message instead of many small messages

Some implementations are slow.

*Current MPI Standard to be found at: <https://www.mpi-forum.org/docs/mpi-4.0/mpi40-report.pdf>*

*Credits to A. Emerson and many other colleagues from CINECA for the slides*